


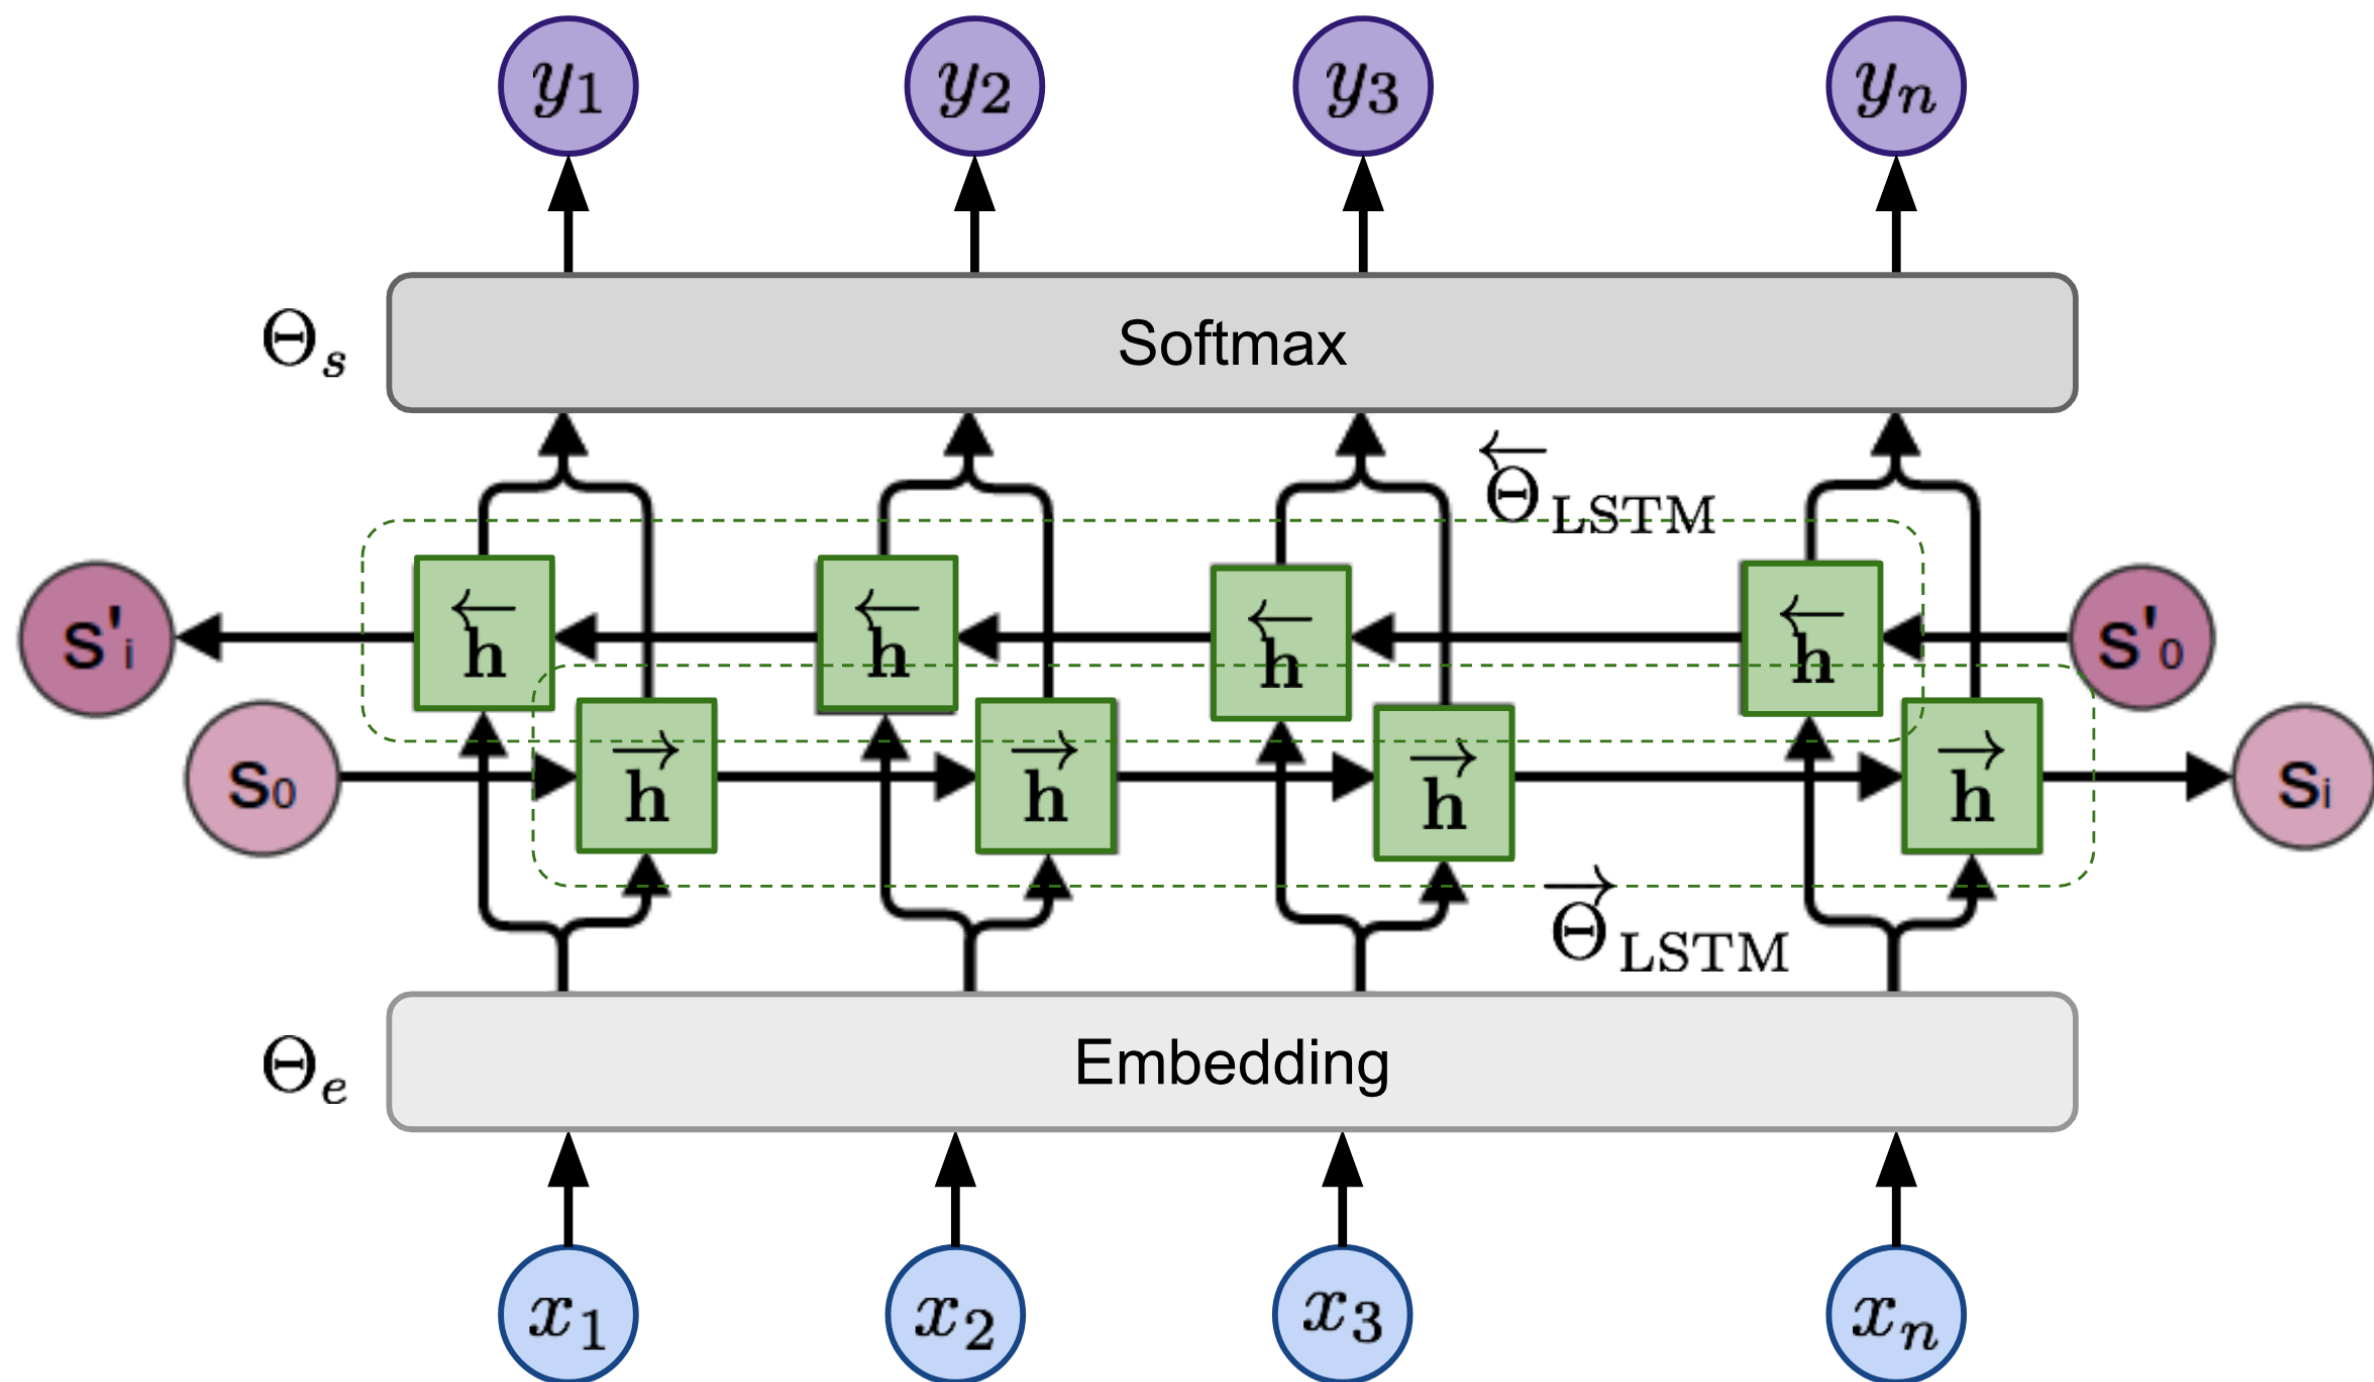
NLP LM Transfer Learning



ELMo

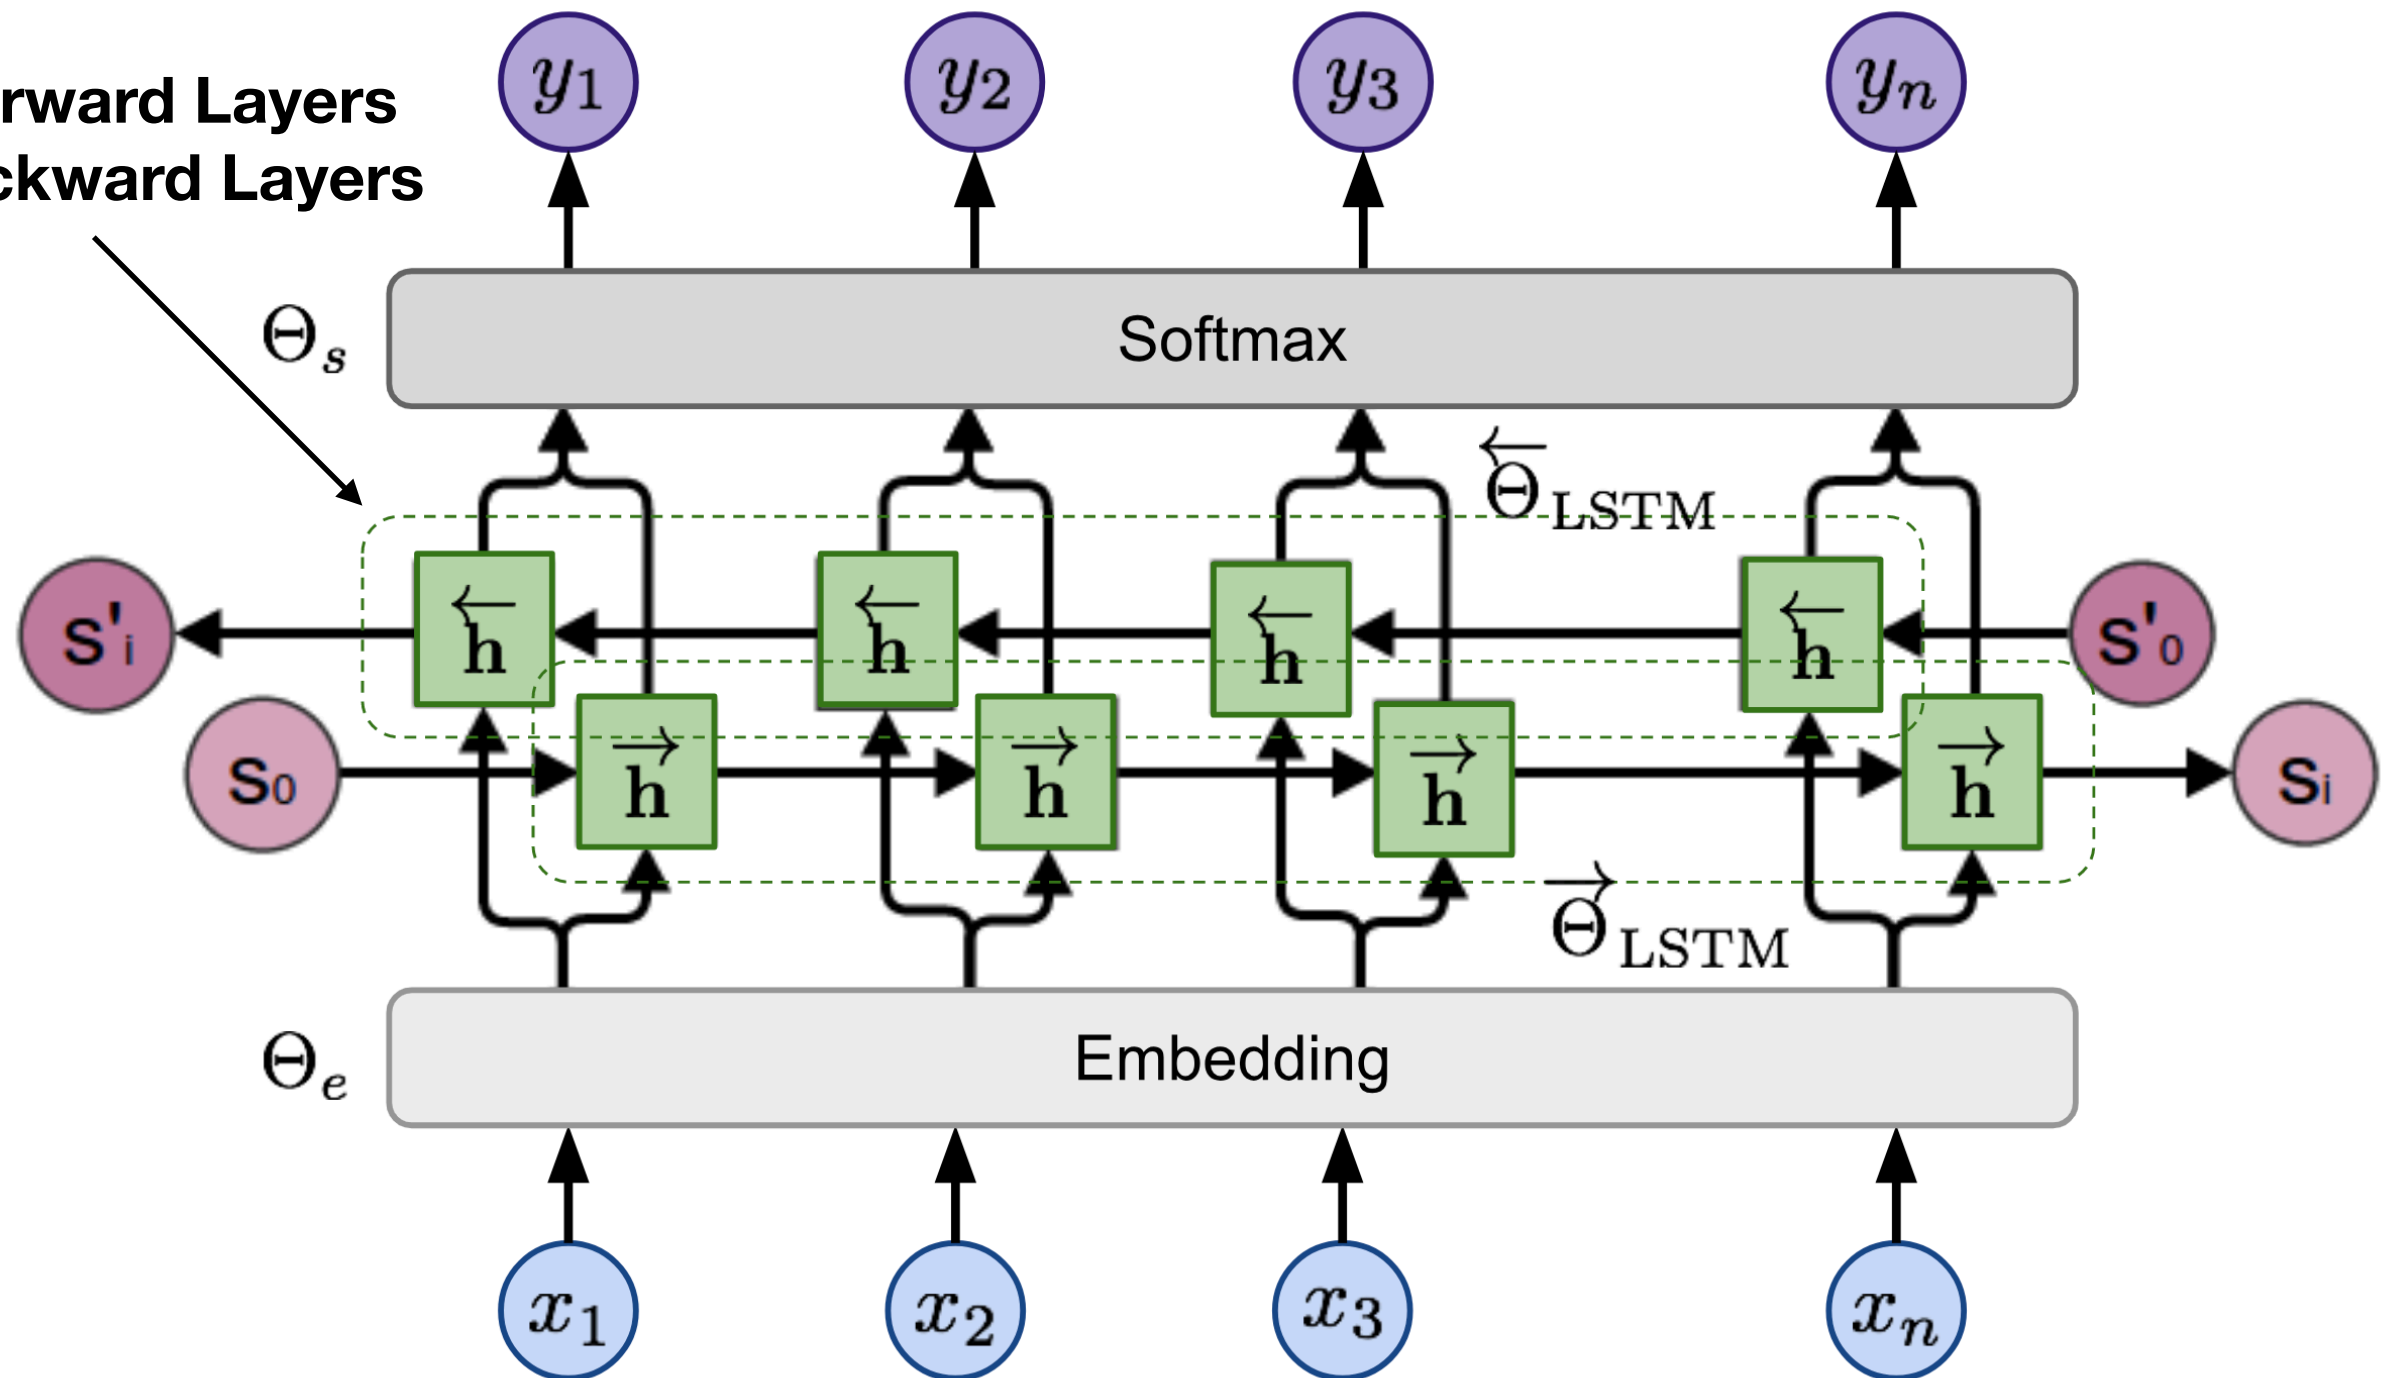


ELMo

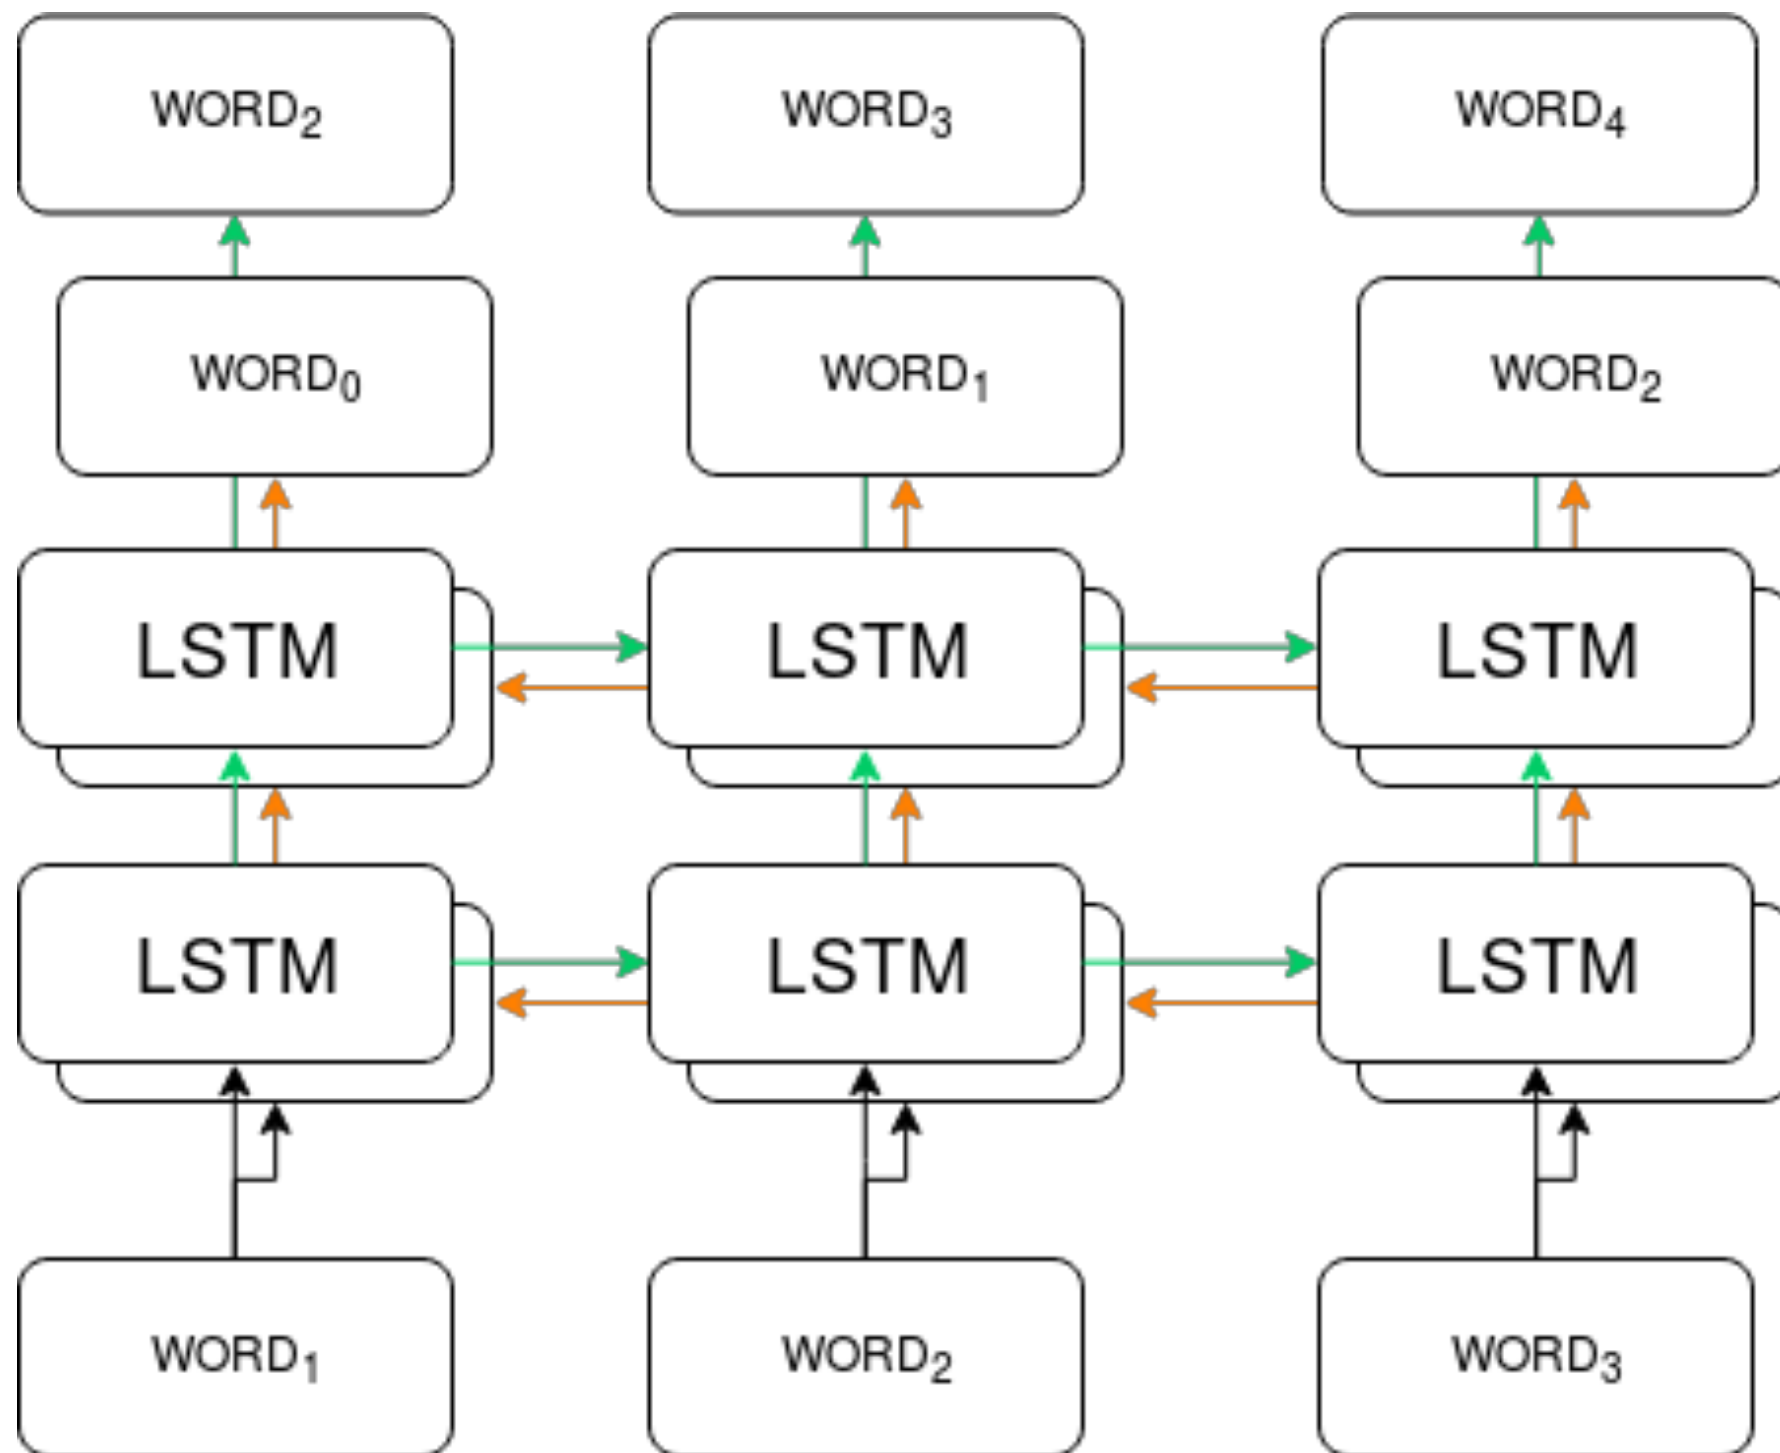


ELMo

2 Forward Layers
2 Backward Layers



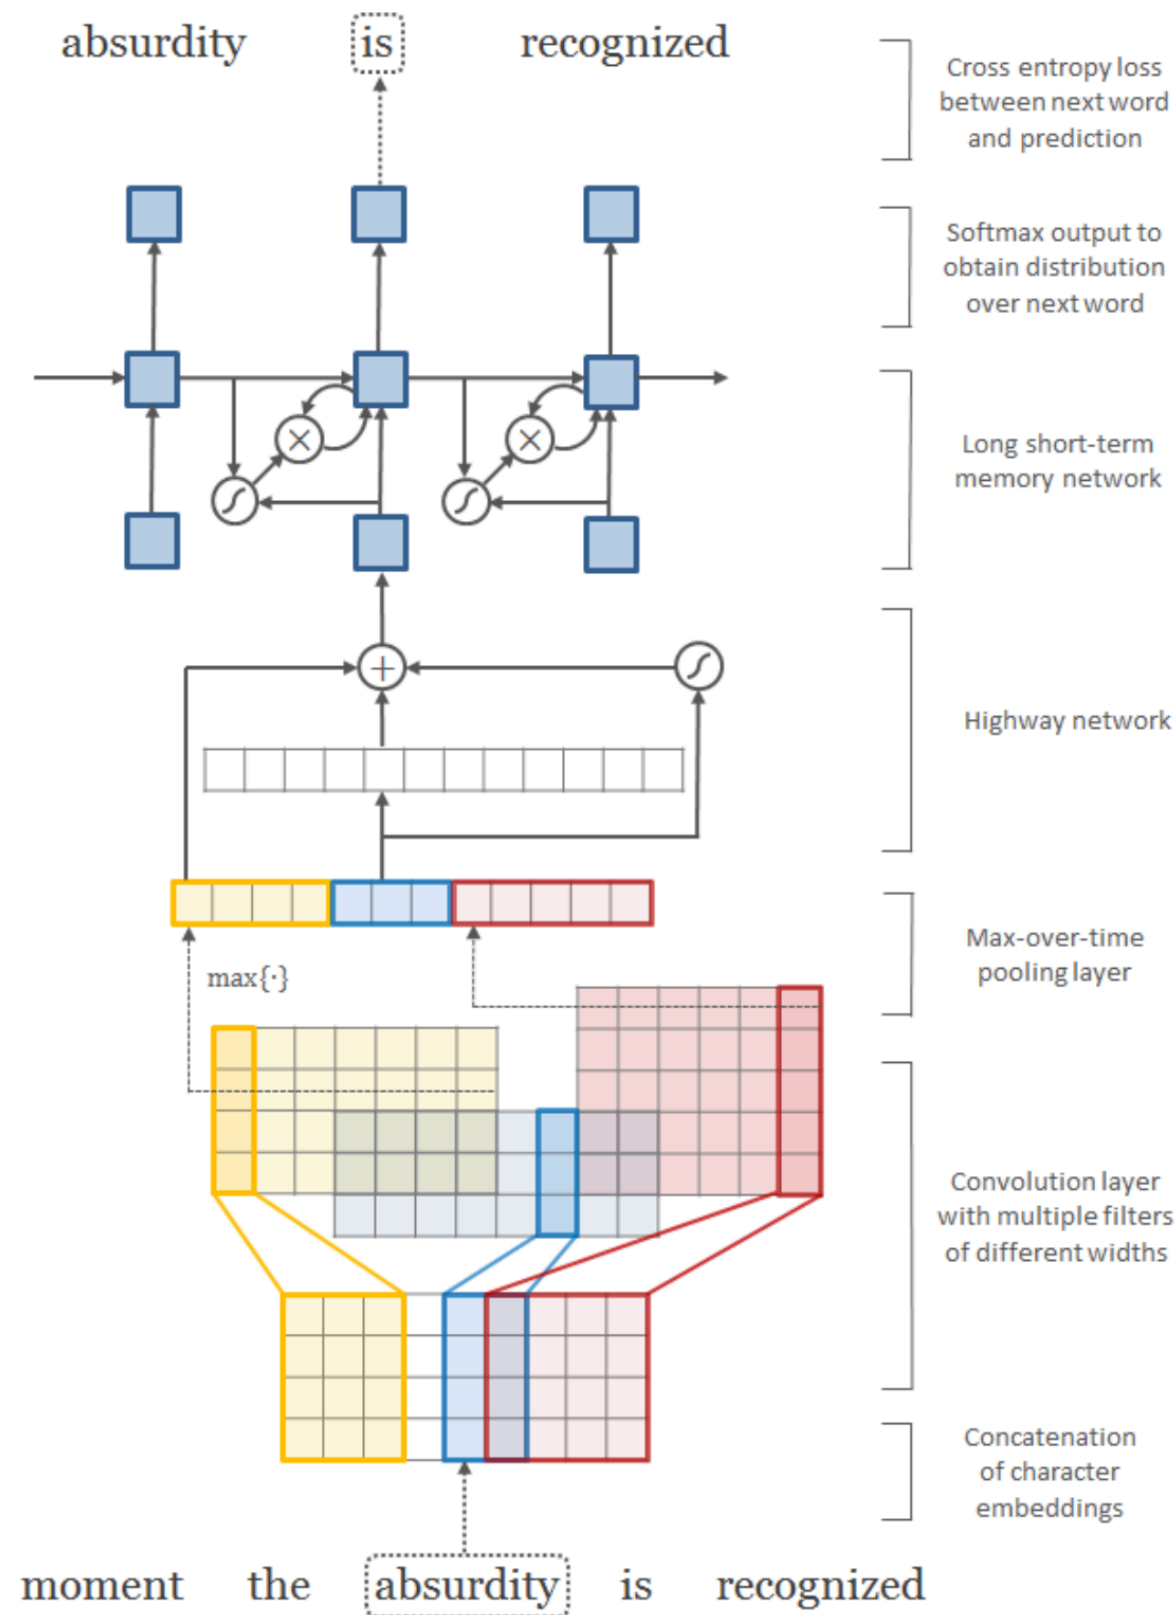
ELMo



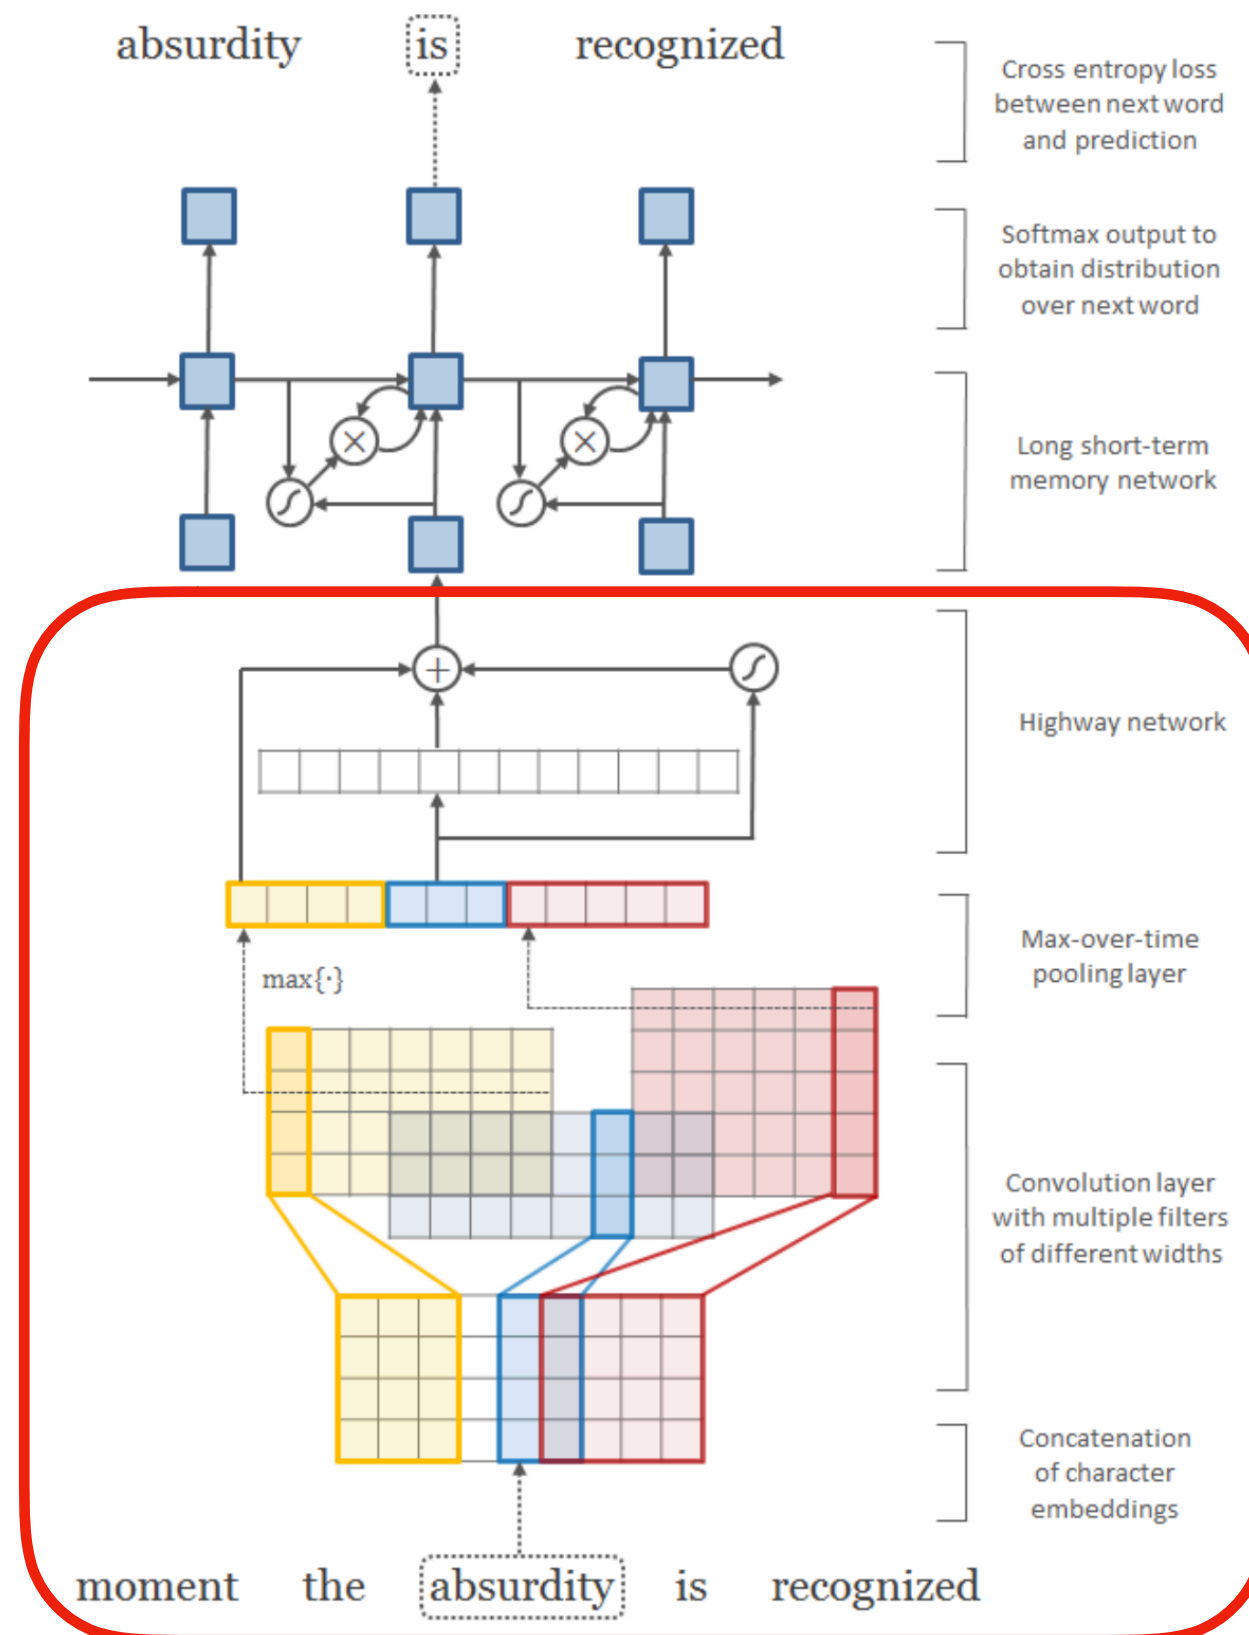
ELMo

Source	<START>	The	poor	don't	have	any	money	<END>
Forward	<START>	The	poor	don't	have	any	money	<END>
Backward	<END>	money	any	have	don't	poor	The	<START>

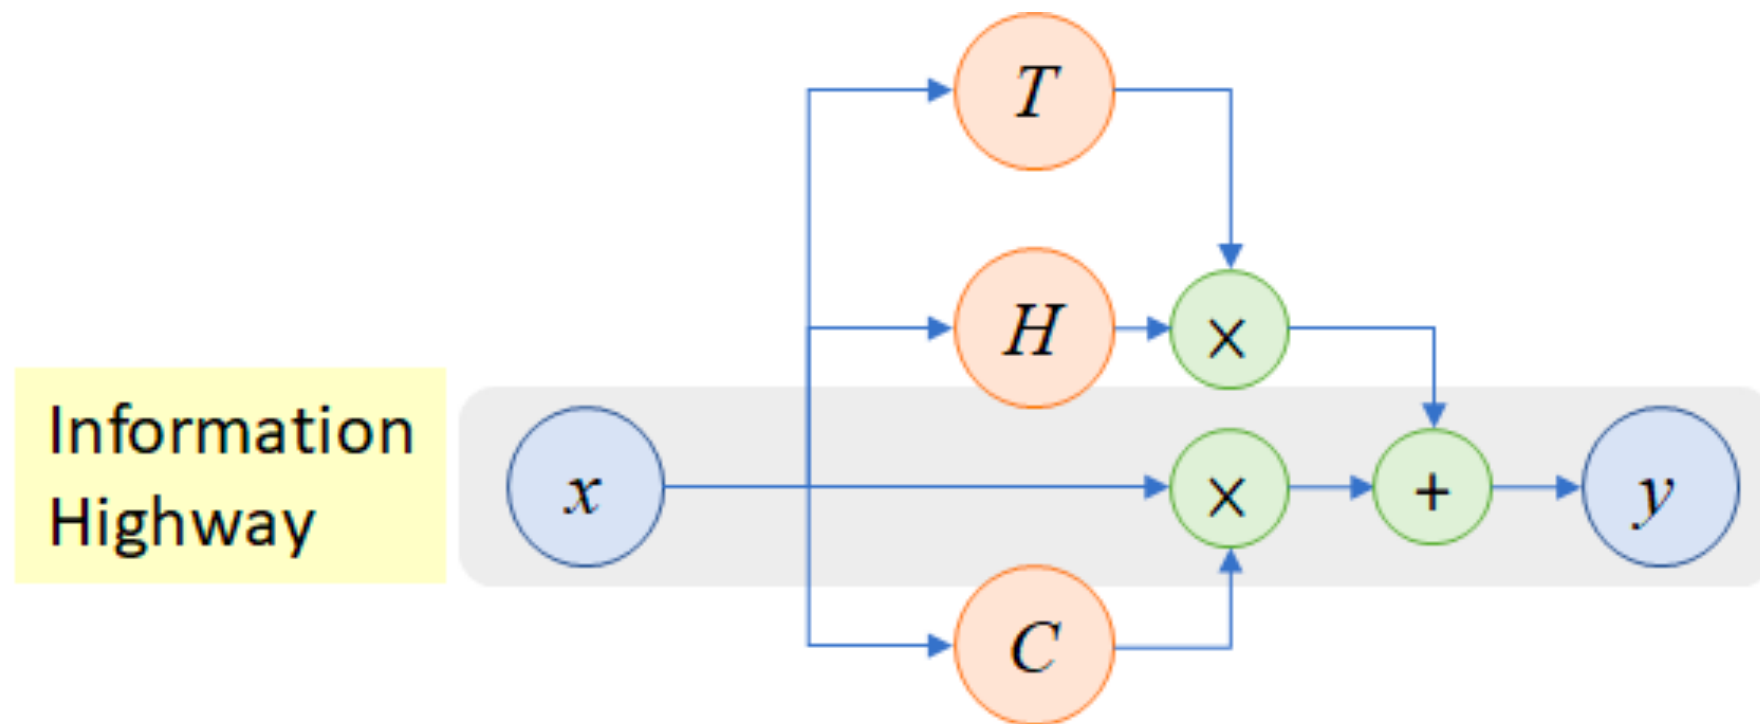
Character-Aware Neural Language Models



Character-Aware Neural Language Models



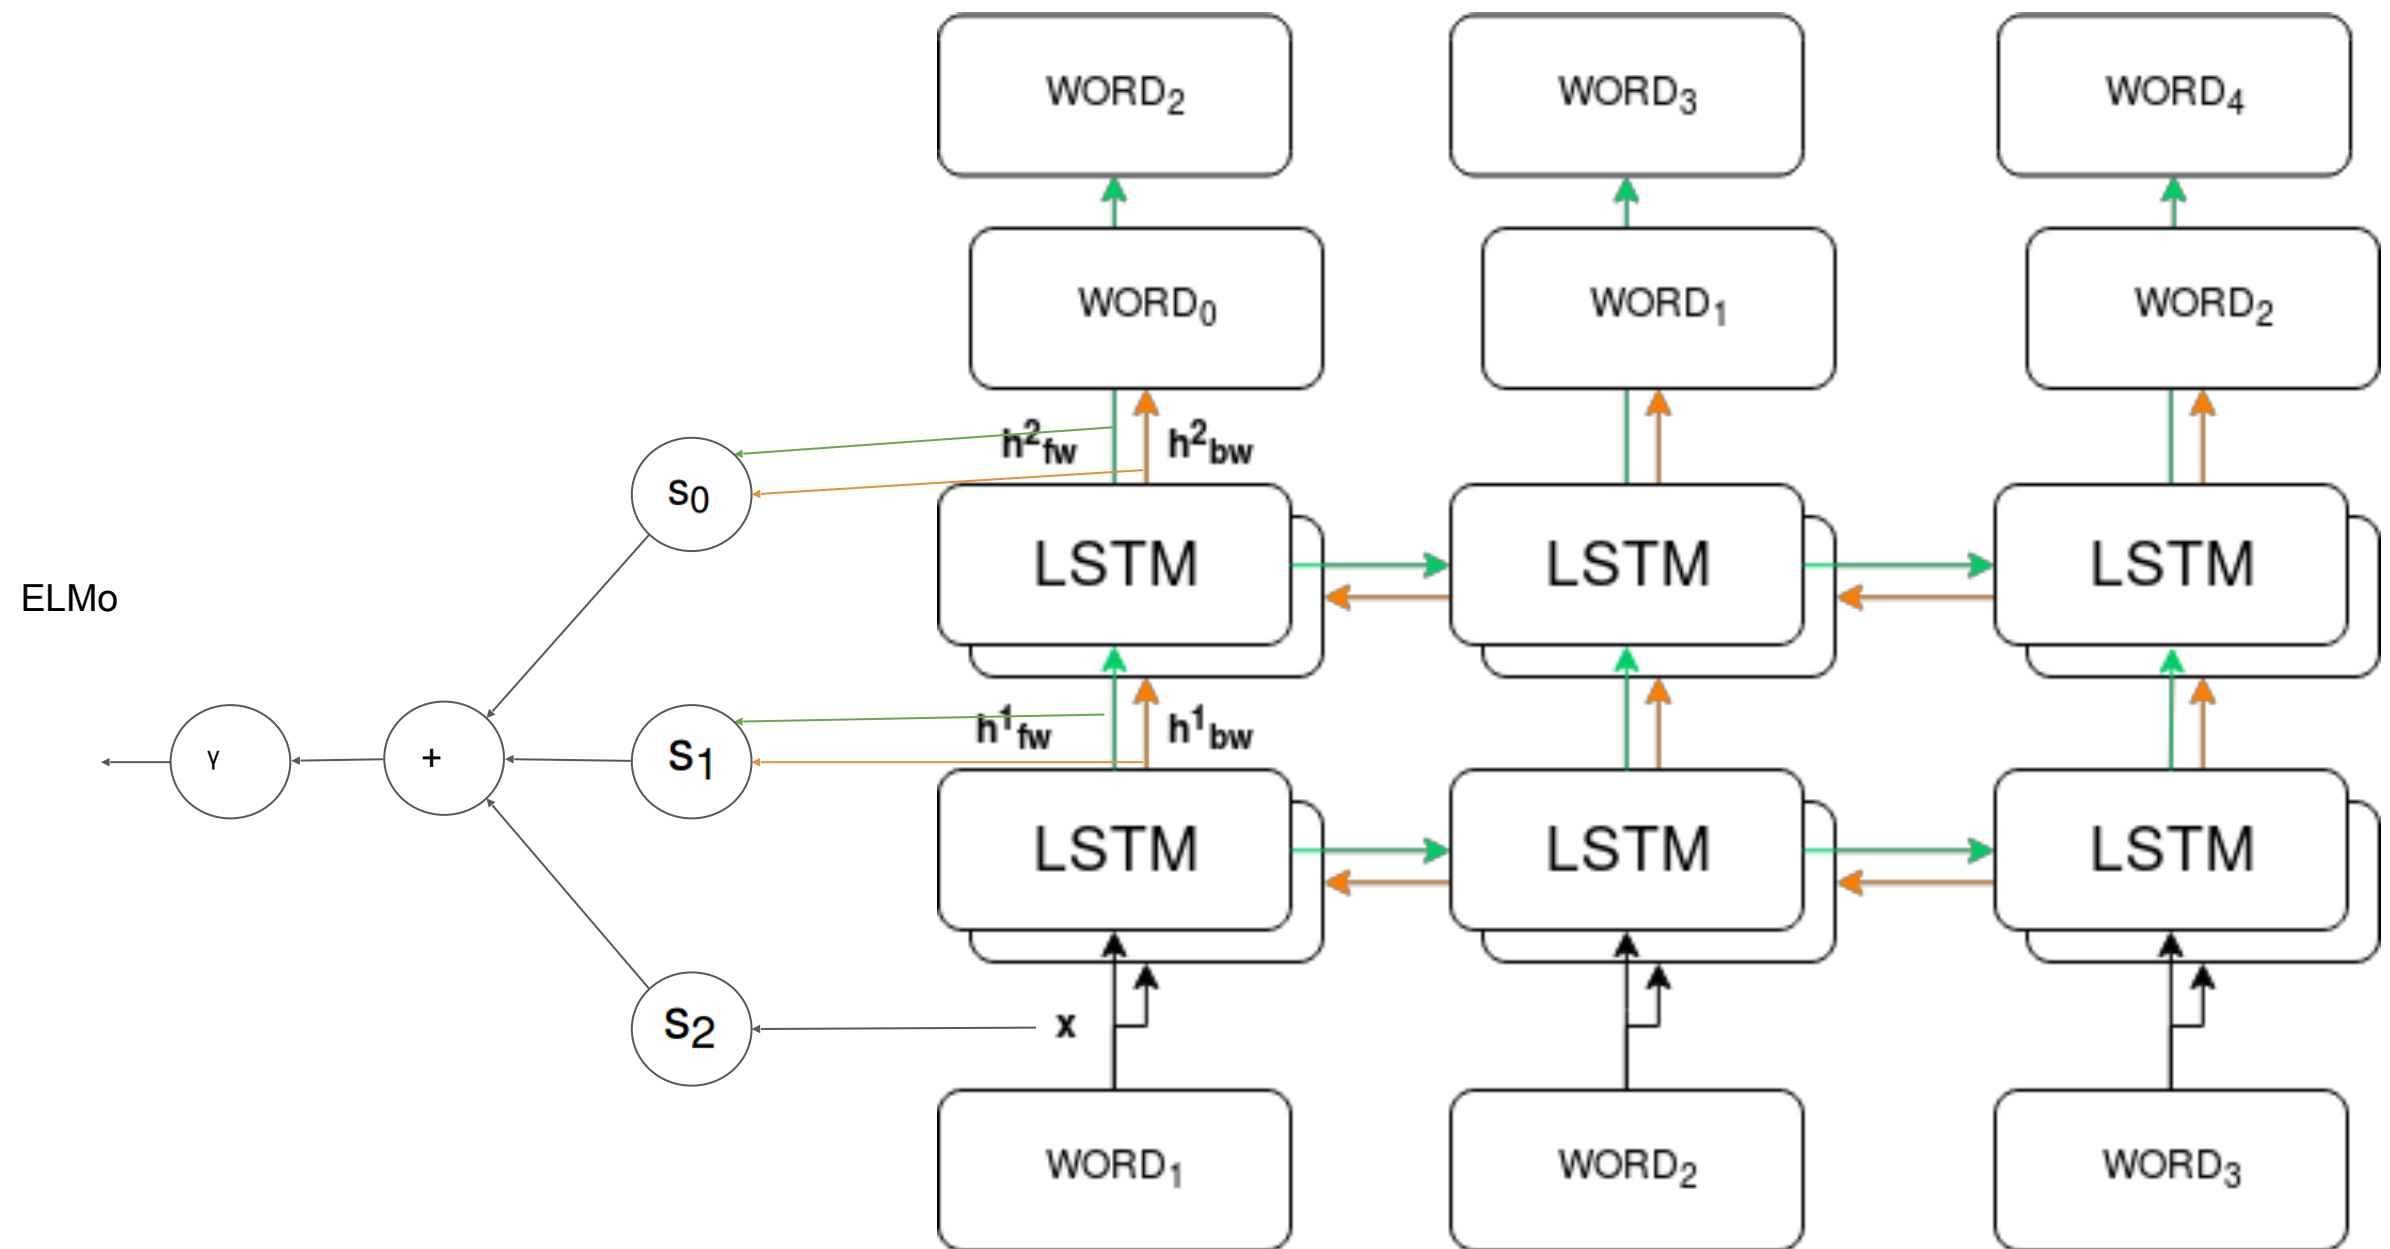
Highway Network



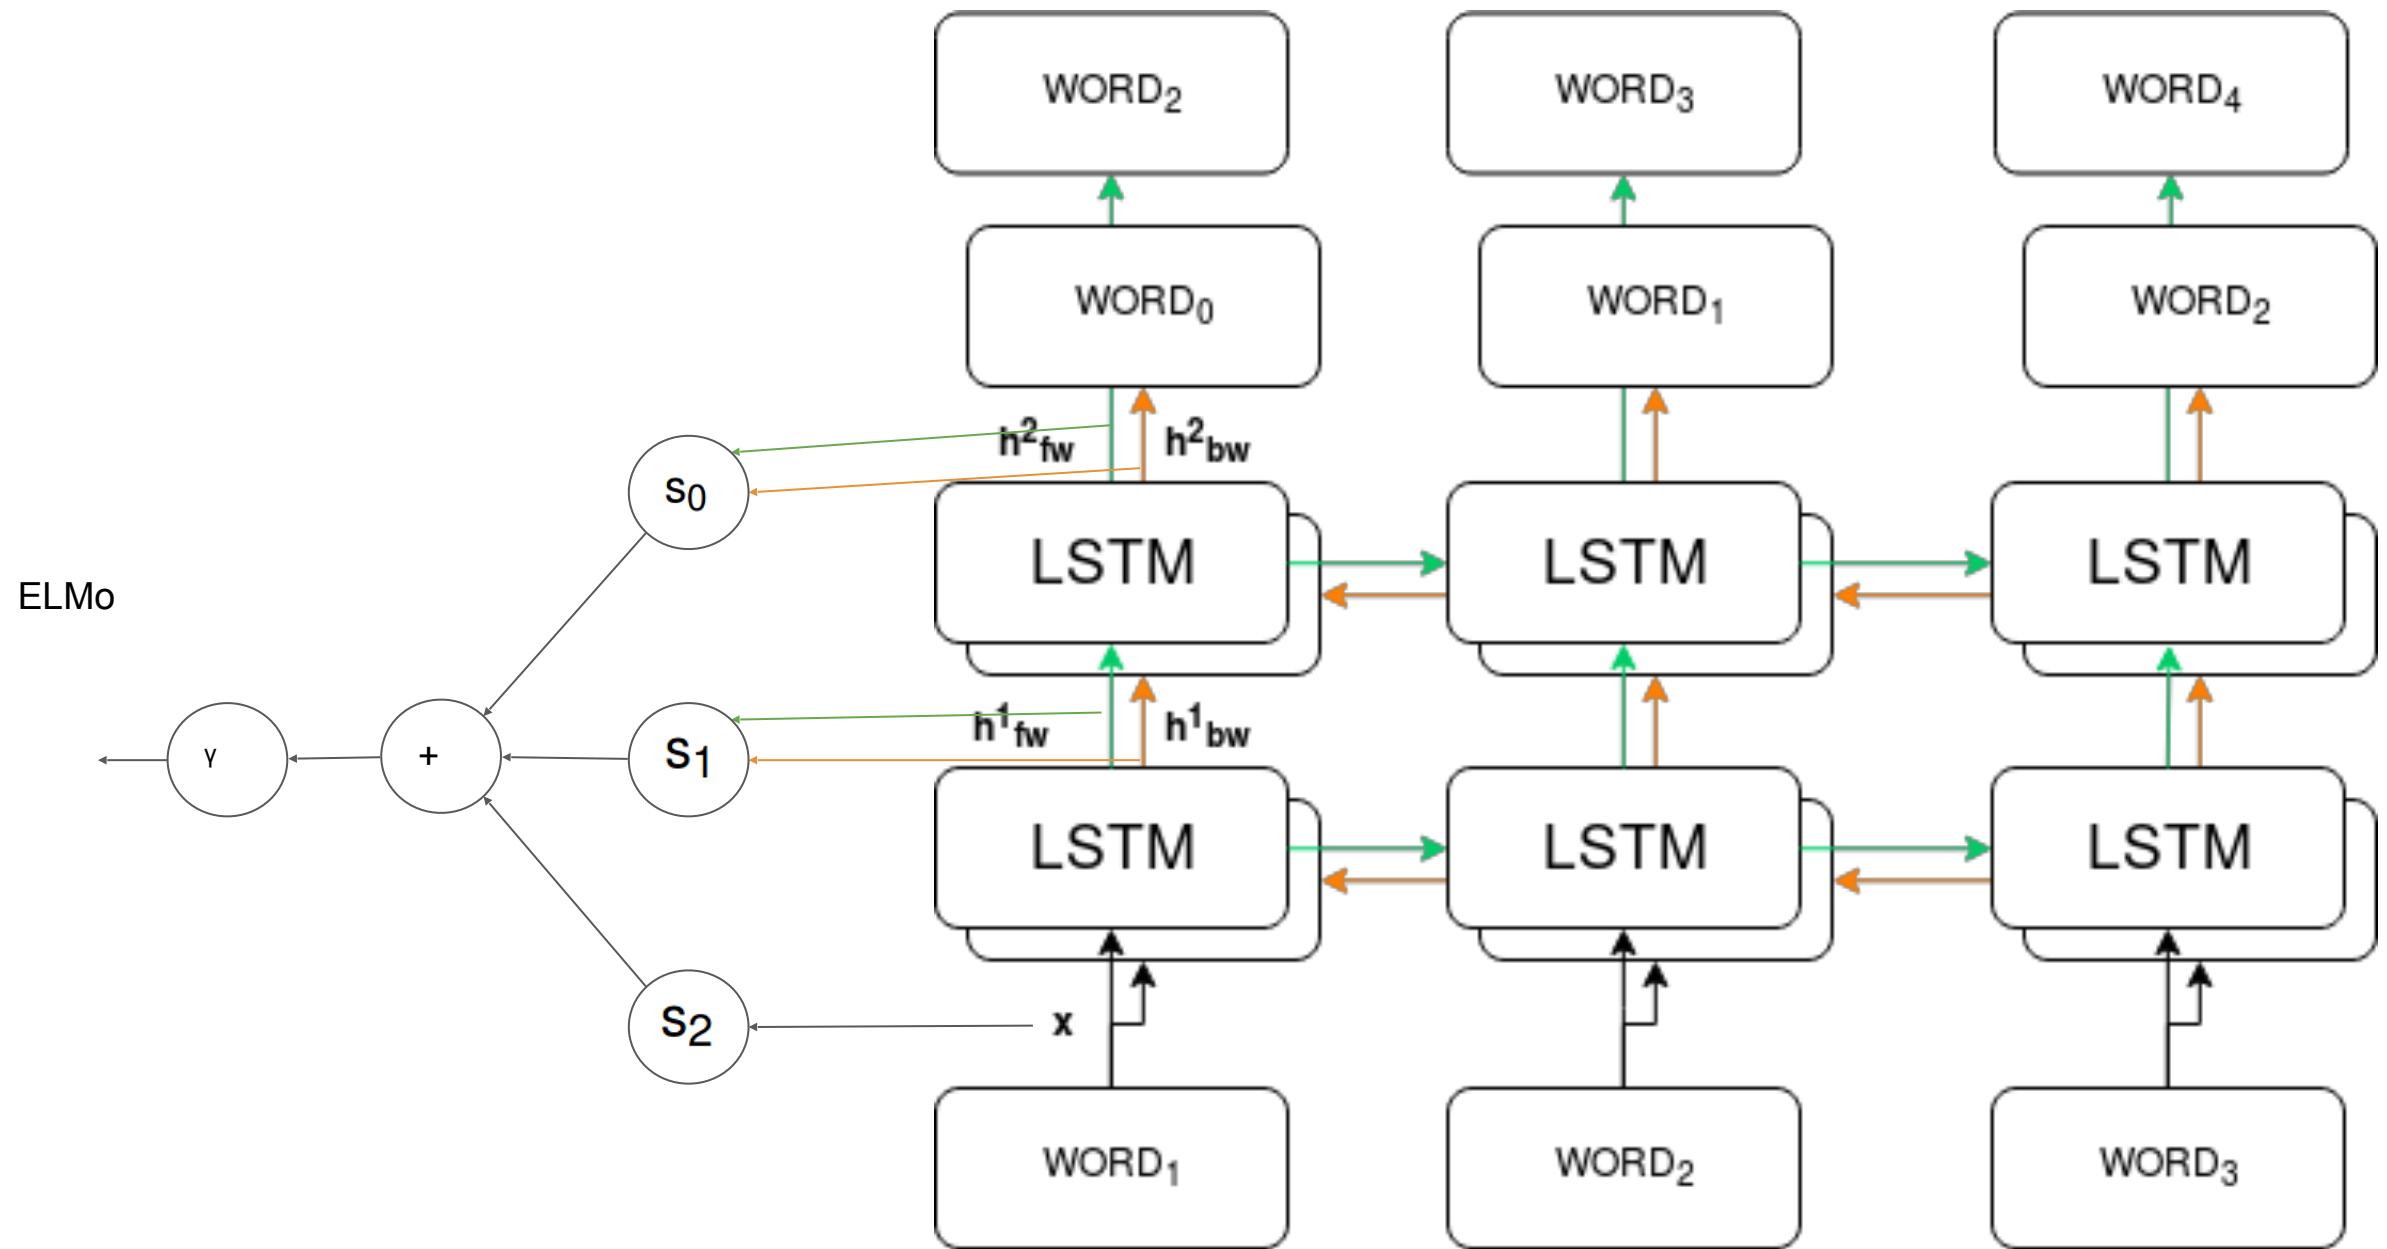
$$y = H(\mathbf{x}, \mathbf{W}_H) \cdot T(\mathbf{x}, \mathbf{W}_T) + \mathbf{x} \cdot C(\mathbf{x}, \mathbf{W}_C).$$

$$y = H(\mathbf{x}, \mathbf{W}_H) \cdot T(\mathbf{x}, \mathbf{W}_T) + \mathbf{x} \cdot (1 - T(\mathbf{x}, \mathbf{W}_T)).$$

ELMo



ELMo



$$\begin{aligned}
 R_k &= \{\mathbf{x}_k^{LM}, \vec{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L\} \\
 &= \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L\},
 \end{aligned}$$

$$\text{ELMo}_k^{\text{task}} = E(R_k; \Theta^{\text{task}}) = \gamma^{\text{task}} \sum_{j=0}^L s_j^{\text{task}} \mathbf{h}_{k,j}^{LM}.$$

ELMo

Results





Figure 2: Visualization of softmax normalized biLM layer weights across tasks and ELMo locations. Normalized weights less than $1/3$ are hatched with horizontal lines and those greater than $2/3$ are speckled.

Tokenization

Закон Ципфа



Tokenization

Word level

- + Small text length
- Big vocabulary size
- OOV

Character level

- Long text
- + Small vocabulary size
- + Almost no OOV

Tokenization

Word level

- + Small text length
- Big vocabulary size
- OOV

Character level

- Long text
- + Small vocabulary size
- + Almost no OOV

1. Word level

i'm a second year student in an ivy league school ->

["i'm", 'a', 'second', 'year', 'student', 'in', 'an', 'ivy', 'league', 'school']

2. Character level

['i', "'", 'm', ' ', 'a', ' ', 's', 'e', 'c', 'o', 'n', 'd', ' ', 'y', 'e', 'a', 'r', ' ', 's', 't', 'u', 'd', 'e', 'n', 't', ' ', 'i', 'n', ' ', 'a', 'n', ' ', 'i', 'v', 'y', ' ', 'l', 'e', 'a', 'g', 'u', 'e', ' ', 's', 'c', 'h', 'o', 'o', 'l']

Tokenization

Word level

- + Small text length
- Big vocabulary size
- OOV

Character level

- Long text
- + Small vocabulary size
- + Almost no OOV



BPE

I saw a girl with a telescope. ->

['__I', '__saw', '__a', '__girl', '__with', '__a', '__', 'te', 'le', 's', 'c', 'o', 'pe', '.']

опубликовано видео убитого саудовского журналиста джамаля хашкуджи ->

['__опубликовано', '__видео', '__убитого', '__саудов', 'ского', '__журналиста',
['__джама', 'ля', '__ха', 'шку', 'джи']

BPE

Algorithm 1 Learn BPE operations

```
import re, collections

def get_stats(vocab):
    pairs = collections.defaultdict(int)
    for word, freq in vocab.items():
        symbols = word.split()
        for i in range(len(symbols)-1):
            pairs[symbols[i], symbols[i+1]] += freq
    return pairs

def merge_vocab(pair, v_in):
    v_out = {}
    bigram = re.escape(' '.join(pair))
    p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
    for word in v_in:
        w_out = p.sub(' '.join(pair), word)
        v_out[w_out] = v_in[word]
    return v_out

vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
        'n e w e s t </w>':6, 'w i d e s t </w>':3}
num_merges = 10
for i in range(num_merges):
    pairs = get_stats(vocab)
    best = max(pairs, key=pairs.get)
    vocab = merge_vocab(best, vocab)
    print(best)
```

r ·	→	r·
l o	→	lo
l o w	→	low
e r·	→	er·

- learning
 - word:freq : {low:5, lowest:2, newer:6, wider:3}
 - marge & count
 1. 'r' '</w>' : 9 → marge'r</w>'
 2. 'e' 'r</w>' : 9 → marge'er</w>'
 3. 'l' 'o' : 7 → marge'lo'
 4. 'lo' 'w' : 7 → marge'low'

→ OOV : 'lower' segmented 'low er</w>'

Vocabulary sizes:

5000, 10000, 15000, ..., 50000

Figure 1: BPE merge operations learned from dictionary {'low', 'lowest', 'newer', 'wider'}.

BPE

SentencePiece

build failing build passing coverage 98% issues 35 open code quality A pypi package 0.1.83 contributions welcome
License Apache 2.0

SentencePiece is an unsupervised text tokenizer and detokenizer mainly for Neural Network-based text generation systems where the vocabulary size is predetermined prior to the neural model training. SentencePiece implements **subword units** (e.g., **byte-pair-encoding (BPE)** [Sennrich et al.]) and **unigram language model** [Kudo.] with the extension of direct training from raw sentences. SentencePiece allows us to make a purely end-to-end system that does not depend on language-specific pre/postprocessing.

This is not an official Google product.

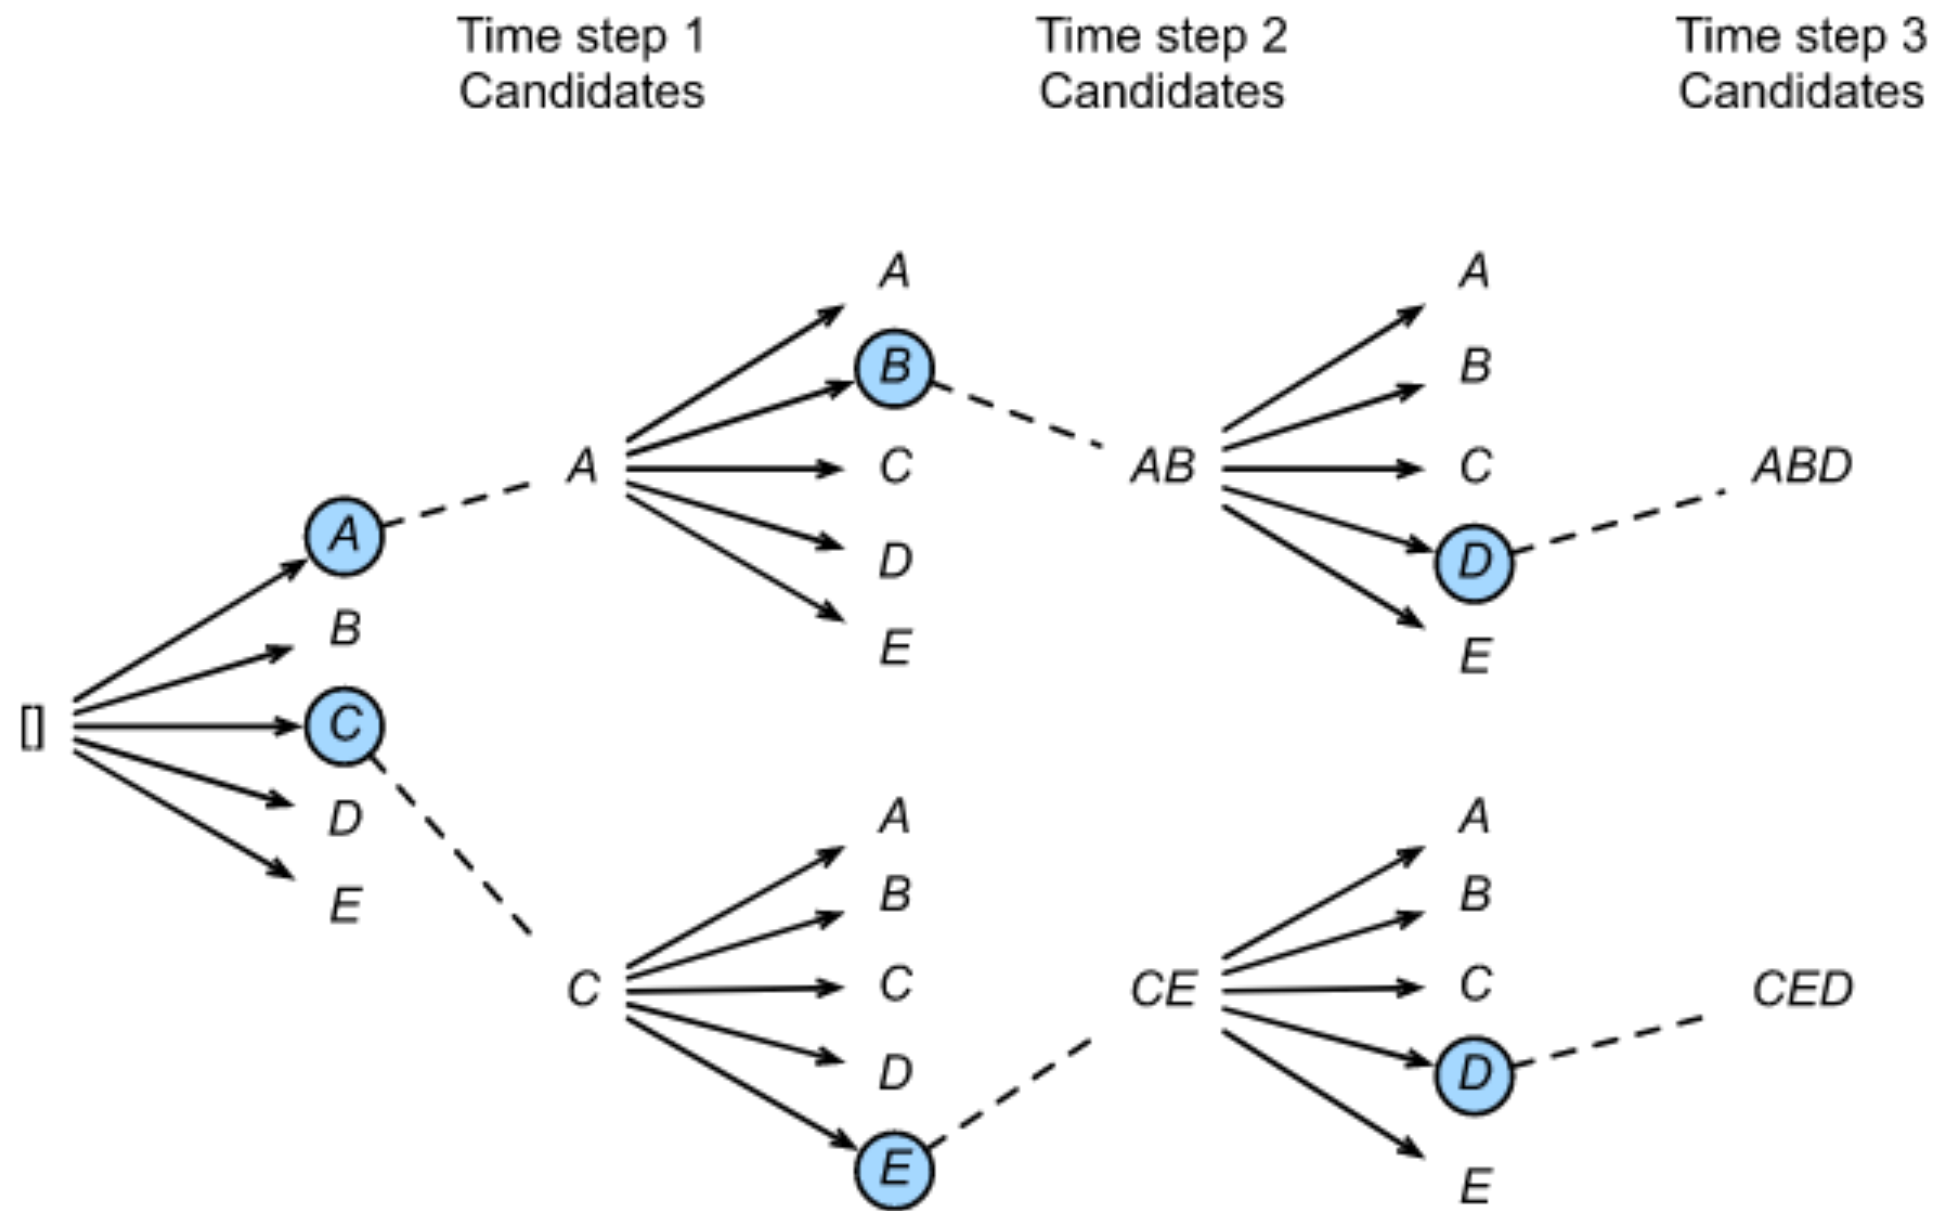
YouTokenToMe

YouTokenToMe is an unsupervised text tokenizer focused on computational efficiency. It currently implements fast Byte Pair Encoding (BPE) [Sennrich et al.]. Our implementation is much faster in training and tokenization than both [fastBPE](#) and [SentencePiece](#). In some test cases, it is 90 times faster. Check out our [benchmark](#) results.

Key advantages:

- Multithreading for training and tokenization
- The algorithm has $O(N)$ complexity, where N is the length of training data
- Highly efficient implementation in C++
- Python wrapper and command-line interface

Beam Search



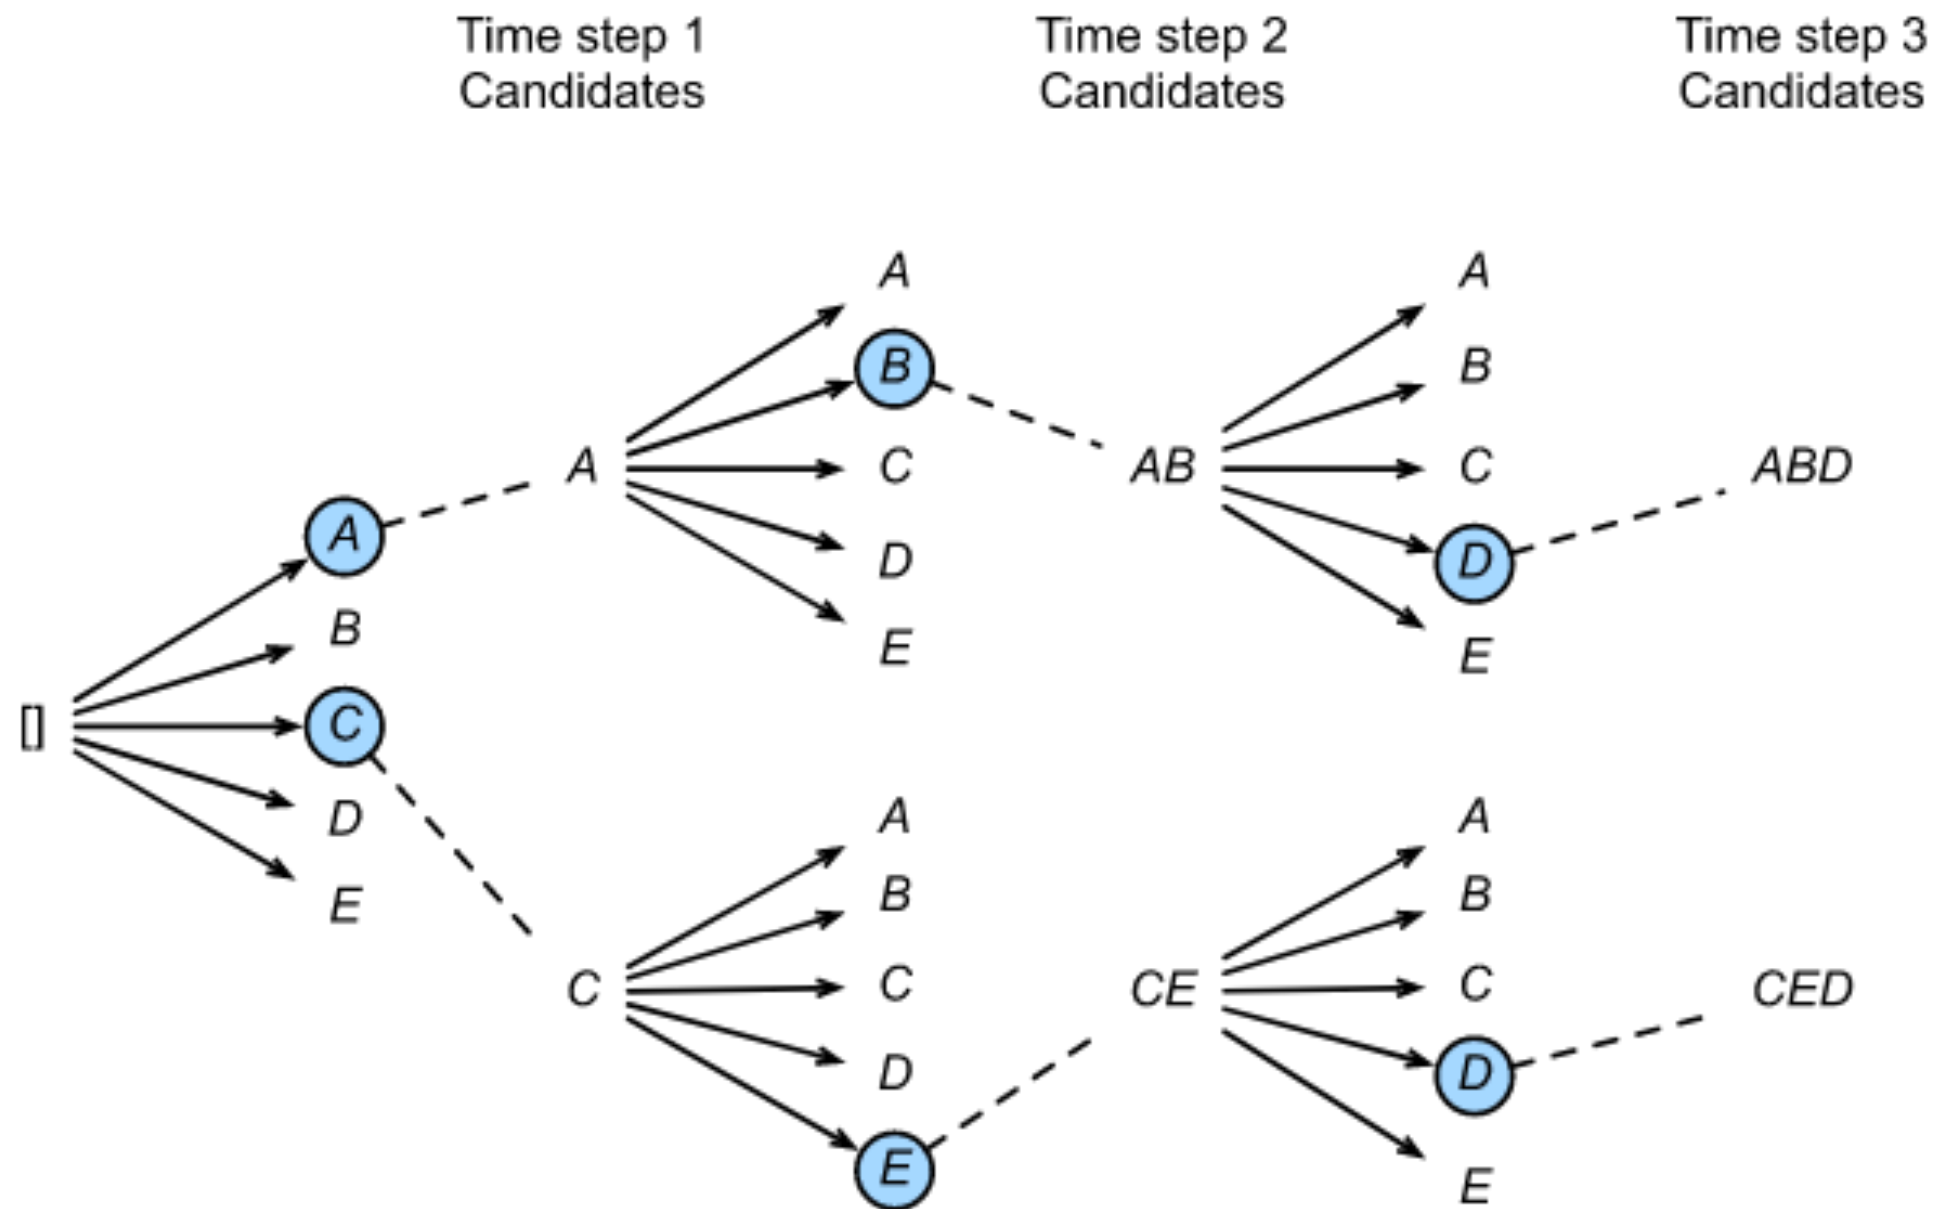
Beam Search

$$p(\mathbf{x}) = \prod_i p(x|x_{<i}) = p(x_0)p(x_1|x_0)p(x_2|x_0, x_1)\dots$$

$$\arg \max_y \prod_{t=1}^{T_y} P(y^{<t>} | x, y^{<1>}, \dots, y^{<t-1>})$$


$$\arg \max_y \sum_{t=1}^{T_y} \log P(y^{<t>} | x, y^{<1>}, \dots, y^{<t-1>})$$

Beam Search




Sequence to sequence


one to one




one to many




many to one



many to many



many to many



Sequence to sequence

Inference

